

BUS 41201 Big Data Midterm

Shri Lekkala

24 April 2024

Setup

```
# Reddit news data
data = read.csv("RedditNews.csv", header = FALSE, skip = 1)
data[,2:3] = data[,1:2]
data[,1] = paste0("RedditNews_",rownames(data))
date<-data[2] # this is the day of the news

subset<-date=="7/1/16" # let's take a look at news headlines on 7/1/16
# data[subset,3] # we have 24 news headlines
```

```
# Read the DJIA data
dj<-read.csv("DJIA.csv")
head(dj) # Open price, highest, lowest and close price
```

```
##           Date      Open      High      Low      Close      Volume Adj.Close
## 1 2016-07-01 17924.24 18002.38 17916.91 17949.37 82160000 17949.37
## 2 2016-06-30 17712.76 17930.61 17711.80 17929.99 133030000 17929.99
## 3 2016-06-29 17456.02 17704.51 17456.02 17694.68 106380000 17694.68
## 4 2016-06-28 17190.51 17409.72 17190.51 17409.72 112190000 17409.72
## 5 2016-06-27 17355.21 17355.21 17063.08 17140.24 138740000 17140.24
## 6 2016-06-24 17946.63 17946.63 17356.34 17400.75 239000000 17400.75
```

```
ndays<-nrow(dj) # 1989 days
```

```
# Read the words
words<-read.csv("WordsFinal.csv",header=F)
words<-words[,1]
head(words)
```

```
## [1] "ab"          "abandon" "abba"     "abbott"   "abc"      "abduct"
```

```
length(words)
```

```
## [1] 5271
```

```
# Read the word-day pairings
doc_word<-read.table("WordFreqFinal.csv",header=F)

# Create a sparse matrix
spm<-sparseMatrix(
  i=doc_word[,1],
  j=doc_word[,2],
  x=doc_word[,3],
  dimnames=list(id=1:ndays,words=words))
dim(spm)
```

```
## [1] 1989 5271
```

```
# We select only words that occur at least 5 times
cols<-apply(spm,2,sum)
index<-apply(spm,2,sum)>5
spm<-spm[,index]

# and words that do not occur every day
index<-apply(spm,2,sum)<ndays
spm<-spm[,index]

dim(spm) # we end up with 3183 words
```

```
## [1] 1989 3183
```



Question 1 Marginal Significance Screening and False Discovery

Setup

```
# *** FDR *** analysis
spm<-spm[-ndays,]
time<-dj[-ndays,1]

# Take returns
par(mfrow=c(1,2))
R<-(dj[-ndays,7]-dj[-1,7])/dj[-1,7]
plot(R~as.Date(time),type="l")

# Take the log of the maximal spread
V<-log(dj[-ndays,3]-dj[-ndays,4])
plot(V~as.Date(time),type="l")
```



```
# FDR: we want to pick a few words that correlate with the outcomes (returns and volatility)

# create a dense matrix of word presence
P <- as.data.frame(as.matrix(spm>0))

# we will practice parallel computing now
margreg <- function(x){
  fit <- lm(Outcome~x)
  sf <- summary(fit)
  return(sf$coef[2,4])
}
```

1.1 Plot the p-values and comment on their distributions (for both outcomes V and R). Is there enough signal to predict prices and volatility?

```

# **** Analysis for Returns ****
cl <- makeCluster(detectCores())

Outcome<-R
clusterExport(cl,"Outcome")

# run the regressions in parallel
mrgpvals_R <- unlist(parLapply(cl,P,margreg))

stopCluster(cl)

# **** Analysis for Volatility ****
cl <- makeCluster(detectCores())

Outcome <- V
clusterExport(cl,"Outcome")

# run the regressions in parallel
mrgpvals_V <- unlist(parLapply(cl,P,margreg))


stopCluster(cl)

# Combine the p-values into a data frame for easy plotting
pvals <- data.frame>Returns = mrgpvals_R, Volatility = mrgpvals_V)

# Plotting the p-values
# Plotting the histograms of p-values


# Histogram for Returns
hist(mrgpvals_R, breaks = 50, main = "Histogram of P-Values for Returns", xlab = "P-Value", ylab = "Frequency")

```



```
# Histogram for Volatility
```

```
hist(mrgpvals_V, breaks = 50, main = "Histogram of P-Values for Volatility", xlab = "P-Value", ylab = "Frequency")
```



The null hypothesis is a uniform distribution of p-values which would mean all p-values having the same frequency.

The histogram of p-values for returns, R, seems to be fairly uniform but seems to be somewhat skewed to the left, which suggests that there is likely to be little signal to predict the returns.

However for volatility, V, the p-values have a clear spike near 0 and seem to be skewed towards the right. The large amount of p-values near 0 suggests there might indeed be a signal for some words to predict volatility.

1.2 What is the alpha value (p-value cutoff) associated with 10% False Discovery Rate? How many words are significant at this level? (Again analyze both outcomes V and R.) What are the advantages and disadvantages of FDR for word selection?

```
## To find the p-value cut off we first order the p values
mrgpvals_R_ordered <- mrgpvals_R[order(mrgpvals_R, decreasing=F)]
mrgpvals_V_ordered <- mrgpvals_V[order(mrgpvals_V, decreasing=F)]

source("fdr.R")

mgr_p_R = length(mrgpvals_R_ordered)
mrg_cutoff_R <- fdr_cut(mrgpvals_R_ordered, q=0.1)
mrg_cutoff_R
```

```
## [1] 1.026222e-05
```

```
mrg_p_V = length(mrgpvals_V_ordered)
mrg_cutoff_V <- fdr_cut(mrgpvals_V_ordered, q=0.1)
mrg_cutoff_V
```

```
## [1] 0.0003571024
```

At the 10% False Discovery rate, the p-value cutoff is:

- 1.026×10^{-5} (4.s.f) for R
- 3.571×10^{-4} (4.s.f) for V

```
# Number of significant words at alpha level 0.1
significant_R <- sum(mrgpvals_R < mrg_cutoff_R)
significant_R
```

```
## [1] 0
```

```
significant_V <- sum(mrgpvals_V < mrg_cutoff_V)
significant_V
```

```
## [1] 11
```

At this FDR level, there are 0 words are significant for R and 11 words significant for V.

-

The main advantages of using FDR analysis for word selection is the ability to handle large datasets and do computations in parallel. So as an initial stage for analyzing which words are important, this method is computationally efficient compared to others. Further, as we are explicitly controlling the expected number of false discoveries, this method allows for a more reliable selection of significant results.

However, the FDR method relies on the assumption that each test is independent. This is not necessarily true for words in headlines as words can occur simultaneously. Further, we completely disregard the structure of groups of words / phrases by treating each word independently. Another disadvantage is that the FDR analysis only selects words based on the magnitude of the association but not their direction, so we cannot use this to select only words with positive or negative association.

1.3 Now, focus only on volatility V. Suppose you just mark the 20 smallest p-values as significant. How many of these discoveries do you expect to be false? Are the p-values independent? Discuss.

Just marking the 20 smallest p-values without using any testing correction like FDR lacks statistical control over the expected number of false discoveries. So we cannot estimate the expected number of false discoveries. However, if we used the discoveries using a 10% FDR from above, then we can expect 10% of the 11 words, i.e. 1.1 of them to be false discoveries.


However, it is unlikely that the p-values are independent as words in language are usually correlated. So some words are likely to appear next to each other, and the presence of one word could mean that another word is more likely to appear. This is true with news headlines which often have common or repeated phrases, which means that some words appear frequently with other words.

So taking into account the lack of independence, we may assume that the expected number of false discoveries is even higher due to dependencies between words not being taken into account. So in conclusion, a more rigorous statistical method should be applied to find significant words rather than simply marking the smallest n as significant.

Question 2 LASSO Variable Selection and Bootstrapping

2.1 Use the LASSO method to come up with a combination of a few words that predict returns R. Pick a lambda and comment on the in-sample R^2 . Is there enough evidence to conclude that headlines predict returns?

```
# First analyze returns
lasso1 <- gamlr(spm, y=R, lambda.min.ratio=1e-3)
plot(lasso1)
```



The top 10 most predictive words from this model are computed below:

```
betas <- drop(coef(lasso1)) # AICc default selection
non_zero_betas = betas[betas!=0][-1] # Exclude intercept and filter out zero betas

# Choose 10 most predictive words in this model
o<-order(non_zero_betas,decreasing=TRUE)
kable(non_zero_betas[o[1:10]])
```

	x
damn	0.0075899
theyll	0.0053115
elect	0.0051119
jamaican	0.0050146
kiss	0.0041769
trawler	0.0041643
paul	0.0038002
rift	0.0027937
grip	0.0024081
cctv	0.0018005

```
# Pick lambda using the AICc
lambda_1 = lasso1$lambda[which.min(AICc(lasso1))]
paste0("The AICc pick for lambda is: ", lambda_1[1])
```

```
## [1] "The AICc pick for lambda is: 0.000631890583514516"
```

```
# Compute the in-sample R2
dev <- lasso1$deviance[which.min(AICc(lasso1))] # deviance #of the AICc selected model
dev0<- lasso1$deviance[1] # null deviance
paste0("The in-sample $R^2$ is: ", 1-dev/dev0)
```

```
## [1] "The in-sample $R^2$ is: 0.07200122438874"
```

The in-sample R^2 is 0.07200122 which suggests there is not much signal. This means that approximately 7.2% of the variance in returns R is explained by the LASSO model based on word frequencies from the headlines. Having a low R^2 score suggests that is not sufficient evidence to conclude that headlines predict returns.

2.2 Repeat the analysis from (2.1) to predict Volatility instead of Returns.

```
# **** LASSO Analysis of volatility **** #
lasso2 <- gamlr(spm, y=V, lambda.min.ratio=1e-3)
betas2 <- drop(coef(lasso2)) # AICc default selection
non_zero_betas2 = betas2[betas2!=0][-1] # Exclude intercept and filter out zero betas

# Choose 10 most predictive words in this model
o <- order(non_zero_betas2,decreasing=TRUE)
kable(non_zero_betas2[o[1:10]])
```

	x
republican	0.1954297
chunk	0.1632192
fusion	0.1455132
govern	0.1123943
shed	0.0870666
august	0.0866411
bailout	0.0786730
barrel	0.0698835
hussein	0.0682012
medvedev	0.0658347

```
# Pick lambda using the AICc
lambda_2 = lasso2$lambda[which.min(AICc(lasso2))]
paste0("The AICc pick for lambda is: ", lambda_2[1])
```

```
## [1] "The AICc pick for lambda is: 0.0260163567306549"
```



```
# Compute the in-sample R2
dev_2 <- lasso2$deviance[which.min(AICc(lasso2))] # deviance #of the AICc selected model
dev0_2<- lasso2$deviance[1] # null deviance
paste0("The in-sample $R^2$ is: ", 1-dev_2/dev0_2)
```

```
## [1] "The in-sample $R^2$ is: 0.161611573416107"
```

So the in-sample R^2 for this model to predict volatility is greater than the previous one for returns, which suggests a stronger signal for prediction.

Next, add one extra predictor, Volatility on a previous day. In other words, fit a LASSO model for predicting today's volatility V_t using word counts and yesterday's volatility V_{t-1} .

$$V_t = a + bV_{t-1} + x'_t\beta + \epsilon_t$$

What is the in-sample R^2 now? Describe the LASSO path and pick the top 10 strongest coefficients. What is the interpretation of the coefficient of word "terrorist" and of V_{t-1} ?

```
# predict future volatility from past volatility
Previous<-log(dj[-1,3]-dj[-1,4]) # remove the last return
spm2<-cbind(Previous,spm) # add the previous return to the model matrix
colnames(spm2)[1]<-"previous" # the first column is the previous volatility
lasso3<- gamlr(spm2, y=V, lambda.min.ratio=1e-3)


# Pick lambda using the AICc
lambda_3 = lasso3$lambda[which.min(AICc(lasso3))]

# Compute the in-sample R2
dev_3 <- lasso3$deviance[which.min(AICc(lasso3))] # deviance #of the AICc selected model
dev0_3<- lasso3$deviance[1] # null deviance
paste0("The in-sample $R^2$ is: ", 1-dev_3/dev0_3)
```

```
## [1] "The in-sample $R^2$ is: 0.331584990647596"
```

So the in-sample R^2 of the third model is much greater than the previous 2 models, and suggests an even stronger signal, so adding the previous day's volatility increased our predictive power.

```
plot(lasso3)
title("LASSO path for third lasso model")
```



From the plot of the coefficients vs lambda, we observe that the model iteratively selects coefficients of different words to become 0 as the value of lambda increases. The dotted vertical line corresponds to the log value of the lambda selected by the AICc. The coefficients all tend to 0 fairly sequentially as lambda decreases, however there is one coefficient that remains non-zero as lambda increases, and only drops to 0 after a significant rise in lambda. This is likely to be the coefficient for the previous day's volatility (V_{t-1}).

```
# Choose top 10 strongest coefficients in this model
betas3 <- drop(coef(lasso3)) # AICc default selection
non_zero_betas3 = betas3[betas3!=0][-1] # Exclude intercept and filter out zero betas

# Choose 10 most predictive words in this model
o <- order(non_zero_betas3,decreasing=TRUE)
kable(non_zero_betas3[o[1:10]])
```

	x
previous	0.4580944
shed	0.1401080
fusion	0.1141525
republican	0.0920512
pioneer	0.0670934
lowest	0.0629167
chunk	0.0581594
ton	0.0491526
william	0.0463384
medvedev	0.0450435

The top 10 strongest coefficients are listed above.

```
# coefficients
coef_terrorist = betas3[names(betas3)=="terrorist"]
paste0("The coefficient of the word terrorist is: ", coef_terrorist)
```

```
## [1] "The coefficient of the word terrorist is: 0.0178389418998624"
```

```
#  $V_{t-1}$ 
coef_prev_v = betas3[names(betas3)=="previous"][1]
paste0("The coefficient of  $V_{t-1}$  is: ", coef_prev_v)
```

```
## [1] "The coefficient of  $V_{t-1}$  is: 0.458094444835972"
```

So the coefficient of the word “terrorist” suggests a small but positive association between the word appearing in headlines and market volatility that day. So the appearance of the word “terrorist” suggest that the market will increase in volatility by ≈ 0.0178 units keeping everything else constant. This may be

The coefficient of the previous day’s volatility is positive and large suggesting that the previous day’s volatility has a strong positive effect on the current day’s volatility. That is, a increase in the yesterday’s volatility by 1 unit leads to an expected increase in today’s volatility by ≈ 0.458 units, keeping other factors fixed. So V_{t-1} is a critical component in forecasting V_t .

2.3 Consider the estimated of lambda selected by AICc in the model from (2.2) (using both words and V_{t-1}). We want to know how variable this estimate is. The starter script has code to bootstrap the sampling distribution for the λ selected by AICc in this regression.

- What is the standard error for the selected λ ?
- Find the 95% CI for λ ?

```
# Bootstrap to obtain s.e. of 1.s.e. chosen lambda

# We apply bootstrap to approximate the sampling distribution of lambda selected by AICc

# export the data to the clusters
cl <- makeCluster(detectCores())

Outcome<-V
clusterExport(cl,"spm2")
clusterExport(cl,"V")


# run 100 bootstrap resample fits
boot_function <- function(ib){
  require(gamlr)
  fit <- gamlr(spm2[ib,],y=V[ib], lambda.min.ratio=1e-3)
  fit$lambda[which.min(AICc(fit))]
}

boots <- 100
n <- nrow(spm2)
resamp <- as.data.frame(
  matrix(sample(1:n,boots*n,replace=TRUE),
    ncol=boots))

lambda_samp <- unlist(parLapply(cl,resamp,boot_function))
stopCluster(cl)

boxplot(lambda_samp)
title("Boxplot of 100 lambdas from bootsrap samples")
```

Boxplot of 100 lambdas from bootstrap samples



The summary statistics for the bootstrap samples are:

```
summary(lambda_samp)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## 0.0004097 0.0006231 0.0007469 0.0007979 0.0009178 0.0020937
```

```
# Original estimate of lambda
```

```
paste0("The original estimated lambda from 2.2 is: ", lambda_3)
```

```
## [1] "The original estimated lambda from 2.2 is: 0.0230003726568119"
```

```
# Compute Standard Error
```

```
se = sqrt(mean((lambda_samp - lambda_3)**2))
paste0("The standard error for the selected lambda is: ", se)
```

```
## [1] "The standard error for the selected lambda is: 0.02220395911209"
```

```
# Compute the 95% confidence interval
```

```
lower = lambda_3 - 1.96 * se
upper = lambda_3 + 1.96 * se
```

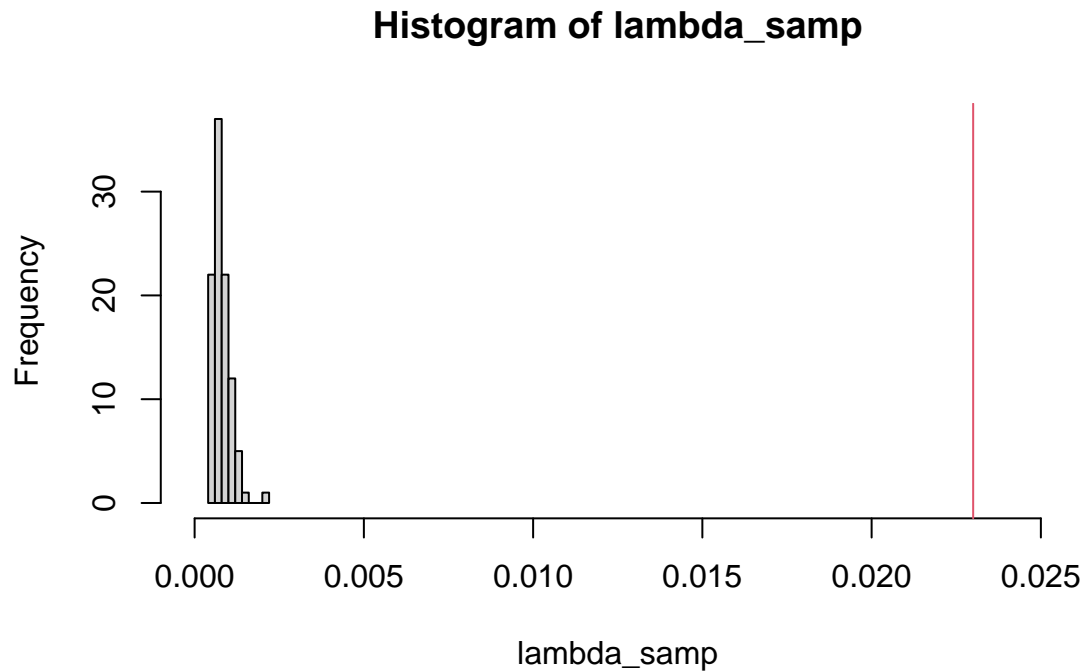
```
print(paste0("The confidence interval for the selected lambda is: [", lower, ",", upper, "]"))
```

```
## [1] "The confidence interval for the selected lambda is: [-0.0205193872028844,0.0665201325165083]"
```

However we note that the original lambda selected in (2.2) is significantly different from the mean of the bootstrapped samples.

So we can repeat the same calculations using the mean of the bootstraps as new estimate for lambda

```
hist(lambda_samp, xlim=c(0, 0.025))  
abline(v=lambda_3, col=2)
```



The red line on the histogram indicates the estimated lambda from the full sample in (2.2). This shows that our estimated lambda is very far from our bootstrap sampling distribution, which suggest

The picture shows that our full sample AICc estimate is near the low end of the sampling distribution; there is high variability in the estimator and this leads us to question the robustness and reliability of the original estimate.

Question 3 High-dimensional Controls and Double LASSO

3.1 Explore a marginal regression (just a regression of V_t on V_{t-1}) to see if there is any correlation. Predict d from x (headlines words), and comment on the degree of confounding we can expect. Is there any information in d independent of x ?

3.2 Isolate the effect of V_{t-1} by running the causal (double) LASSO. Interpret this effect and compare it to the effect obtained from the naive LASSO.

3.3 Can we safely claim that the effect is causal?

Bonus Freestyle Analysis

Provide additional analysis of the data. Points will handed out only for insightful use of data mining tools, not for scattershot application of techniques.